


PLSC 502 – Fall 2022

Estimation and Estimators

October 13, 2022

Remember This?



Random Variables, Take Two

For a random variable X :

$$X_i = \underbrace{\mu}_{\text{"systematic part"}} + \underbrace{u_i}_{\text{"stochastic part"}}$$

where μ is the population mean (expected value) of X and $\text{Cov}(\mu, u) = 0$.

That implies that:

$$\underbrace{u_i}_{\text{"error"}} = \underbrace{X_i}_{\text{"observed"}} - \underbrace{\mu}_{\text{"expected"}}$$

Random Variables, Take Two

What's our expectation for u ?

$$\begin{aligned}E(u) &= E(X - \mu) \\&= E(X) - E(\mu) \\&= E(X) - \mu \\&= \mu - \mu \\&= 0\end{aligned}$$

and so:

$$\begin{aligned}\text{Var}(X) &= E[(X - \mu)^2] \\&= E(u^2)\end{aligned}$$

and

$$\begin{aligned}\text{Var}(u) &= E[(u - E(u))^2] \\&= E[(u - 0)^2] \\&= E(u^2).\end{aligned}$$

Estimation Example: \bar{X}

Challenge: Estimate $\mu = E(X)$ from a sample of N observations.

$$\begin{aligned}\bar{X} &= \frac{1}{N} \sum_{i=1}^N X_i \\ &= \frac{1}{N} \sum_{i=1}^N (\mu + u_i) \\ &= \frac{1}{N} \sum_{i=1}^N (\mu) + \frac{1}{N} \sum_{i=1}^N (u_i) \\ &= \frac{1}{N} (N\mu) + \frac{1}{N} \sum_{i=1}^N (u_i) \\ &= \mu + \bar{u}\end{aligned}$$

The point: \bar{X} is a random variable.

Small-Sample Properties

- Hold irrespective of N
- “Small sample estimators”

Large-Sample (Asymptotic) Properties

- Hold as $N \rightarrow \infty$
- “More is better”

Unbiasedness

Start with a generic population parameter θ , and an estimator of it $\hat{\theta}$ based on a sample of N observations...

Unbiasedness means:

$$E(\hat{\theta}) = \theta$$

“Bias” is:

$$B(\hat{\theta}) = E(\hat{\theta}) - \theta$$

Example: For \bar{X} , we know that:

$$\begin{aligned} E(\bar{X}) &= E(\mu + \bar{u}) \\ &= E(\mu) + E(\bar{u}) \\ &= \mu + 0 \\ &= \mu \end{aligned}$$

and so:

$$B(\bar{X}) = 0.$$

Multiple Unbiased Estimators

For $N = 2$:

$$Z = \lambda_1 X_1 + \lambda_2 X_2.$$

note that

$$\begin{aligned} E(Z) &= E(\lambda_1 X_1 + \lambda_2 X_2) \\ &= E(\lambda_1 X_1) + E(\lambda_2 X_2) \\ &= \lambda_1 E(X_1) + \lambda_2 E(X_2) \\ &= \lambda_1 \mu + \lambda_2 \mu \\ &= (\lambda_1 + \lambda_2) \mu \end{aligned}$$

Means

$$E(Z) = \mu \iff (\lambda_1 + \lambda_2) = 1.0$$

and in fact:

$$E(Z) = \mu \iff \sum_{i=1}^N \lambda_i = 1.0.$$

Q: Why do we use $\lambda_i = \frac{1}{N} \forall i$?

Efficiency:


- is *relative variability* – how much difference we would expect in our $\hat{\theta}$ s from one sample to the next...
- ...so a more efficient estimator has higher “reliability.”
- ...is related to **information** (specifically, the *Fisher information* in the sample).

Note that:

- To be *fully efficient*¹, an estimator must be unbiased. BUT...
- ...the least-variance estimator need not be an unbiased one.

¹That is, to achieve the *Cramer-Rao lower bound*, something we'll discuss in detail a bit later.

Efficiency: Unbiased $\hat{\theta}$ s



Efficiency (continued)

Note that for our example with $N = 2$, where $\text{Var}(X) = \sigma^2$:

$$\begin{aligned}\text{Var}(Z) &= \text{Var}(\lambda_1 X_1 + \lambda_2 X_2) \\ &= (\lambda_1^2 + \lambda_2^2)\sigma^2\end{aligned}$$

and:

$$\begin{aligned}\lambda_1^2 + \lambda_2^2 &= \lambda_1^2 + (1 - \lambda_1)^2 \\ &= \lambda_1^2 + (1 - 2\lambda_1 + \lambda_1^2) \\ &= 2\lambda_1^2 - 2\lambda_1 + 1.\end{aligned}$$

Minimize!

$$\begin{aligned}\frac{\partial 2\lambda_1^2 - 2\lambda_1 + 1}{\partial \lambda_1} &= 4\lambda_1 - 2 \\ 4\lambda_1 - 2 &= 0 \\ \lambda_1 &= 0.5\end{aligned}$$

Mean Squared Error

The “mean squared error” (“MSE”) of an estimator $\hat{\theta}$ is:

$$\begin{aligned}\text{MSE}(\hat{\theta}) &= E[(\hat{\theta} - \theta)^2] \\ &= E[B(\hat{\theta})^2] \\ &= \text{Var}(\hat{\theta}) + [B(\hat{\theta})]^2\end{aligned}$$

Note that:

- The MSE of an unbiased estimator is equal to its variance [that is, $\text{MSE} = \text{Var}(\hat{\theta})$].
- Among unbiased estimators, the efficient estimator will always have the smallest MSE [because $B(\hat{\theta}) = [B(\hat{\theta})]^2 = 0$].

Comparing Estimators via MSE

As an estimator of μ , \bar{X} has:

- $B(\bar{X}) = 0$
- $\text{Var}(\bar{X}) = \sigma^2/N$, so
- $\text{MSE}(\bar{X}) = \sigma^2/N + (0)^2 = \sigma^2/N$.

My alternative: the “Six Estimator”!

$$\zeta = 6$$

(That's a “zeta.” Gotta learn your Greek letters.)

Comparing Estimators via MSE


Properties of ζ (for $\zeta = 6$):

$$\begin{aligned}B(\zeta) &= E(\zeta - \mu) \\&= E(6) - E(\mu) \\&= 6 - \mu,\end{aligned}$$

$$\begin{aligned}\text{Var}(\zeta) &= \text{Var}(6) \\&= 0\end{aligned}$$

and so:

$$\begin{aligned}\text{MSE}(\zeta) &= \text{Var}(\zeta) + [B(\zeta)]^2 \\&= 0 + (6 - \mu)^2 \\&= 36 - 12\mu + \mu^2\end{aligned}$$



The black line is the MSE of ζ , expressed as a function of the “true” population mean μ . The other colored lines are the MSEs for \bar{X} , under the assumption that $\sigma^2 = 10$ and $N = \{20, 100, 1000\}$, respectively.

Large-Sample Properties: Consistency

An estimator $\hat{\theta}$ is *consistent* if:

$$\lim_{N \rightarrow \infty} \Pr[|\hat{\theta} - \theta| < \epsilon] = 1.0$$


for an arbitrarily small $\epsilon > 0$

Equivalently:

$$E(\hat{\theta}_N) \rightarrow \theta \text{ as } N \rightarrow \infty$$

Intuition: “Asymptotic unbiasedness” ...

A Consistent Estimator $\hat{\theta}$



Among estimators:

- Unbiased $>$ Consistent $>$ Biased
- Fully Efficient $>$ Asymptotically Efficient $>$ Inefficient
- MSE is one way to trade off bias vs. efficiency

Estimation Example: The Poisson

Recall the *Poisson* distribution:

$$f(x) \equiv \Pr(X = x) = \frac{\lambda^x \exp(-\lambda)}{x!}.$$


for $x \in \{0, 1, 2, \dots\}$.

The Poisson:


- ...is a distribution for *counts* of *independent events*;
- ...is a *one parameter* distribution, where
- ...the parameter λ is both the *mean* and the *variance* of X .

Poisson Densities


Lambda = 0.5




Lambda = 1



Lambda = 4



Lambda = 8



Poisson Estimation

What is a “good” estimator for λ ?

For a series of N i.i.d. values $\{X_1, X_2, \dots, X_N\}$ drawn from a Poisson distribution, their *joint* probability is:

$$f(X_1, X_2, \dots, X_N | \lambda) \equiv f(\mathbf{X}) = \prod_{i=1}^N \frac{\lambda^{X_i} \exp(-\lambda)}{X_i!}. \quad (1)$$

This is sometimes known as the *likelihood* (more on that later...), and it relies on the fact that the joint probability of two independent random variables equals the product of the two marginal probabilities:

$$\Pr(A, B \mid A \perp B) = \Pr(A) \times \Pr(B)$$

Poisson Estimation

We can simplify (1) by taking its log:

$$\begin{aligned}\ln[f(\mathbf{X})] &= \ln \left[\prod_{i=1}^N \frac{\lambda^{X_i} \exp(-\lambda)}{X_i!} \right] \\ &= \sum_{i=1}^N \ln \left[\frac{\lambda^{X_i} \exp(-\lambda)}{X_i!} \right] \\ &= \sum_{i=1}^N [X_i \ln(\lambda) - \lambda - \ln(X_i!)] \\ &= -N\lambda + \ln(\lambda) \sum_{i=1}^N X_i - \sum_{i=1}^N \ln(X_i!)\end{aligned}$$

(This is the *log-likelihood*...)

Poisson Estimation

If we want to know the value of λ that maximizes this joint (log-)probability, we can figure that out too:

$$\frac{\partial \ln f(\mathbf{X})}{\partial \lambda} = -N + \frac{1}{\lambda} \sum_{i=1}^N X_i$$

and then:

$$-N + \frac{1}{\lambda} \sum_{i=1}^N X_i = 0$$

and so:

$$\hat{\lambda} = \frac{1}{N} \sum_{i=1}^N X_i$$

IOW, one version of a “good” estimator for λ (the “maximum likelihood estimator”) is the empirical mean \bar{X} ...

Poisson Mean Characteristics

What can we say about this $\hat{\lambda}$?

$$\begin{aligned}E(\hat{\lambda}) &= E\left[\frac{1}{N}\sum_{i=1}^N X_i\right] \\&= \frac{1}{N}\sum_{i=1}^N E(X_i) \\&= \frac{1}{N}\sum_{i=1}^N \lambda \\&= \lambda\end{aligned}$$

so:

$$B(\hat{\lambda}) = 0 \text{ (unbiasedness)}$$

Also: Because $\text{Var}(X) = \lambda$, this also means that $\hat{\lambda}$ is also an unbiased estimate of the variance.

More Poisson Mean Characteristics

Variance / efficiency?

Because $\hat{\lambda}$ is unbiased, we know that:

$$\text{MSE}(\hat{\lambda}) = \text{Var}(\hat{\lambda}).$$

Central limit theorem means that:

$$\hat{\lambda} \sim N\left(\lambda, \frac{\lambda}{N}\right)$$

so:

$$\text{MSE}(\hat{\lambda}) = \frac{\lambda}{N}.$$

Example One: Simulation

The Plan:

1. Draw N values of X from a Poisson distribution with a known value of λ ;
2. Calculate $\hat{\lambda} = \bar{X}$;
3. Repeat steps (1) - (2) many times;
4. Examine the distribution of the $\hat{\lambda}$ s

Details


- Vary $\lambda \in \{0.2, 1.0, 8.0\}$
- Vary $N \in \{5, 50, 500\}$

A Little Code


```
> L<-c(0.2,1,8) # the lambdas:
> N<-c(5,50,500) # the Ns:
> sims<-4000      # number of sims
> Out<-data.frame(matrix(nrow=sims,ncol=length(N)*length(L)))
>
> c <- 0           # column indicator for "Out"
> set.seed(7222009) # Seed
>
> for(i in 1:length(N)) { # Looping over sample sizes...
+   for(j in 1:length(L)) { # Looping over lambdas
+     c <- c+1             # increment column indicator
+     for(k in 1:sims) {   # Looping over 4000 simulations each
+       df<-rpois(N[i],L[j]) # Draw N values from Poisson(lambda)
+       Out[k,c]<-mean(df)   # Store the mean of the N draws
+       rm(df)
+     }
+   }
+ }
```

A Picture


Lambda=0.2, N=5




Lambda=1, N=5




Lambda=8, N=5




Lambda=0.2, N=50




Lambda=1, N=50




Lambda=8, N=50




Lambda=0.2, N=500



Lambda=1, N=500



Lambda=8, N=500



Example Two: "Real" Data

Back to the English Premier League!

> PL

	Rank	Team	GamesPlayed	Won	Drew	Lost	GoalsFor	GoalsAgainst	GoalDifference	Points
2	1	Arsenal	9	8	0	1	23	10	13	24
3	2	Manchester City	9	7	2	0	33	9	24	23
4	3	Tottenham Hotspur	9	6	2	1	20	10	10	20
5	4	Chelsea	8	5	1	2	13	10	3	16
6	5	Manchester United	8	5	0	3	13	15	-2	15
7	6	Newcastle United	9	3	5	1	17	9	8	14
8	7	Brighton and Hove Albion	8	4	2	2	14	9	5	14
9	8	Bournemouth	9	3	3	3	8	20	-12	12
10	9	Fulham	9	3	2	4	14	18	-4	11
11	10	Liverpool	8	2	4	2	20	12	8	10
12	11	Brentford	9	2	4	3	16	17	-1	10
13	12	Everton	9	2	4	3	8	9	-1	10
14	13	West Ham United	9	3	1	5	8	10	-2	10
15	14	Leeds United	8	2	3	3	11	12	-1	9
16	15	Crystal Palace	8	2	3	3	10	12	-2	9
17	16	Aston Villa	9	2	3	4	7	11	-4	9
18	17	Southampton	9	2	1	6	8	17	-9	7
19	18	Wolverhampton Wanderers	9	1	3	5	3	12	-9	6
20	19	Nottingham Forest	9	1	2	6	7	22	-15	5
21	20	Leicester City	9	1	1	7	15	24	-9	4

Premier League: Summary

```
> describe(PL)
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
Rank*	1	20	10.5	5.92	10.5	10.50	7.41	1	20	19	0.00	-1.38	1.32
Team*	2	20	10.5	5.92	10.5	10.50	7.41	1	20	19	0.00	-1.38	1.32
GamesPlayed	3	20	8.7	0.47	9.0	8.75	0.00	8	9	1	-0.81	-1.41	0.11
Won	4	20	3.2	2.02	2.5	2.94	0.74	1	8	7	0.95	-0.25	0.45
Drew	5	20	2.3	1.38	2.0	2.31	1.48	0	5	5	0.05	-0.98	0.31
Lost	6	20	3.2	1.88	3.0	3.12	1.48	0	7	7	0.31	-0.86	0.42
GoalsFor	7	20	13.4	6.92	13.0	12.62	7.41	3	33	30	1.01	0.89	1.55
GoalsAgainst	8	20	13.4	4.68	12.0	12.75	4.45	9	24	15	0.87	-0.62	1.05
GoalDifference	9	20	0.0	9.36	-1.5	-0.62	10.38	-15	24	39	0.65	0.06	2.09
Points	10	20	11.9	5.52	10.0	11.38	5.19	4	24	20	0.75	-0.34	1.24

Fitting a Poisson Distribution

```
> library(MASS)
> PoisMean <- fitdistr(PL$Drew,"poisson")
> PoisMean
      lambda
      2.3000000
(0.3391165)
>
> coef(PoisMean)
lambda
      2.3
>
> vcov(PoisMean)
      lambda
lambda  0.115
>
> Note:
>
> coef(PoisMean) / nrow(PL)
lambda
      0.115
```

Actual vs. Theoretical Draws

